

Fine-tuning LLM for Multilingual Code Generation

Link to Google Doc to add comment and see latest version: [click here](#)

OVERVIEW

- **Very** Large code LLMs (> 1B params) often have good performance in a variety of languages ([Salesforce/codegen-2B-multi](#) -> C, C++, Go, Java, JavaScript, and Python)
- **Small** code LLMs (< 1B params) are usually trained to work on one programming language to have good performance ([Salesforce/codegen-350M-mono](#) -> Python)
- Small code LLMs can be extended to cover more languages by a variety of methods:
 - a. Full Fine-tuning:
 - Pros: Best performance in fine-tuned model
 - Cons: High computational requirements, catastrophic forgetting
 - b. LoRa Fine-tuning:
 - Pros: Efficacy, Scalable, No catastrophic forgetting
 - Cons: Lower performance than full fine-tuning
 - c. Knowledge Distillation: tbd

GOALS

1. Optimally extend Small code LLMs to new programming language considering:
 - a. Performance: in terms of Perplexity and HumanEval
 - b. Efficiency: in terms of Memory and Time
2. Study the effect of different factors on the training, evaluation and inference of models:
 - a. Training Data: Selection of github repos and how to filter them
 - b. Training Strategies: learning rate, batch size, gradient accumulation
 - c. Evaluation: Stride and stop tokens
 - d. Inference: Sampling strategy and temperature
3. Qualitative Analysis of the results:
 - a. Failure modes comparison in HumanEval between Python and Java (Exceptions, runtime errors and not passing test cases)
 - b. Correlation between Perplexity and HumanEval

-
4. Further objectives:
 - a. Relation between Different Programming Languages
 - b. Knowledge Distillation

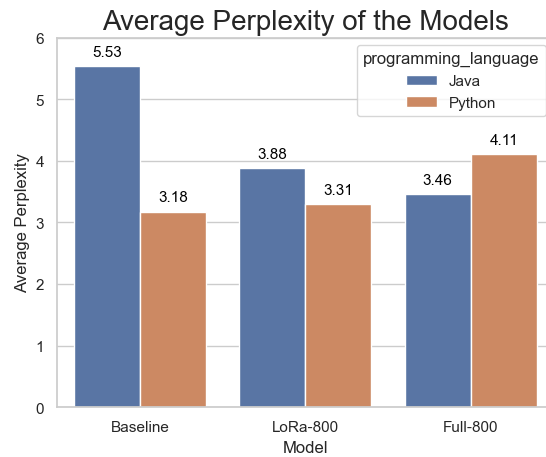
Timeline:

1. Read papers on Code LLM and select appropriate baseline -> [Salesforce/codegen-350M-mono](#)
2. Read papers on Code datasets and select perplexity evaluation data -> [bigcode/the-stack-dedup](#)
3. Read papers on evaluation strategy to implement HumanEval in python and Java -> follow standards of [CodeX](#) and [Multipl-E](#)
4. Setting up Evaluation:
 - o Done: Perplexity Evaluation on Google-Colab and HumanEval on Google-Colab and Virtual Containers (for Java)
 - o Challenges: Perplexity Dataset, HumanEval Stop tokens for java
5. Generate Baseline Results using baseline model
6. Setting up Training and Evaluation environments:
 - o Done: Configure Full & LoRa fine-tuning on Google-Colab, Track training progress on WandB, Save checkpoints on HuggingFace.
 - o Challenges: Resuming Training reset some variables(LR, number of steps), Memory to 15GB and duration to 3 hours
7. Generate Results for the full tuned model and LoRa tuned model.
8. Edit the training setup (learning rate , gradient accumulation and batch size) based on the literature ([bigcode/starcoder](#)) and the results of the first finetuning
9. Train the model again and notice the effect of modification on two models:
 - o Full fine tuning: Significant improvement in Perplexity and HumanEval for **Java** and significant drop for **Python**
 - o LoRa fine tuning: Significant improvement in Perplexity and Slight in HumanEval for **Java** and no change for **Python**
10. Plan Next steps: further modifications, larger models, inference and distillation
11. Setting up live inference environment:
 - o Done: Live inference configured on HuggingFace spaces to generate code from tuned models base on given prompt -> [Code Inference](#)
 - o Challenges: CPU limitation

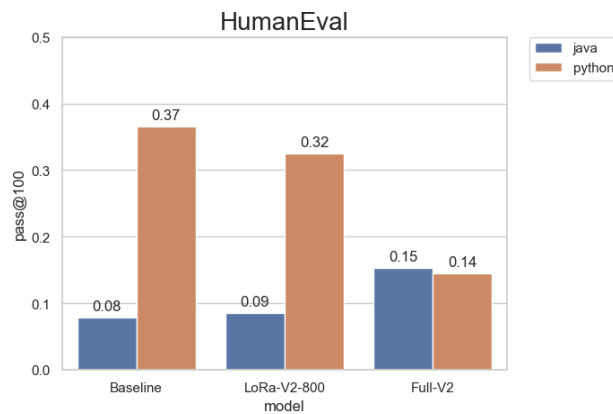
Results

1. Comparisons

- Comparing the models in terms of:
 - i. Perplexity: The Lower the better



- ii. HumanEval: The Higher the better



- iii. Average Memory Usage: given the same setup for all hyperparameters
 - iv. Training Time